

```
In [2]: import os
current_directory = os.getcwd()
print("Current working directory:", current_directory)
```

Current working directory: C:\Users\Prashant Sharma\PycharmProjects

Exploratory Data Analysis Starter

Import packages

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Shows plots in jupyter notebook
%matplotlib inline

# Set plot style
sns.set(color_codes=True)
```


Loading data with Pandas

```
In [6]: client_df = pd.read_csv('client_data.csv')
price_df = pd.read_csv('price_data.csv')
client_df.head(3)
```

```
Out[6]:
```

	id	channel_sales	cons_12m	cons_gas_12m	co
0	24011ae4ebbe3035111d65fa7c15bc57	foosdfpfkusacimwkcsosbicdxkicaua	0	54946	
1	d29c2c54acc38ff3c0614d0a653813dd	MISSING	4660	0	
2	764c75f661154dac3a6c254cd082ea7d	foosdfpfkusacimwkcsosbicdxkicaua	544	0	

3 rows × 26 columns



```
In [7]: price_df.head(3)
```

```
Out[7]:
```

	id	price_date	price_off_peak_var	price_peak_var	price_mid_peak
0	038af19179925da21a25619c5a24b745	2015-01-01	0.151367	0.0	
1	038af19179925da21a25619c5a24b745	2015-02-01	0.151367	0.0	
2	038af19179925da21a25619c5a24b745	2015-03-01	0.151367	0.0	



Descriptive statistics of data

Data types It is useful to first understand the data that you're dealing with along with the data types of each column. The data types may dictate how you transform and engineer features.

To get an overview of the data types within a data frame, use the `info()` method.

In [9]:

```
price_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193002 entries, 0 to 193001
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    193002 non-null object
1   price_date            193002 non-null object
2   price_off_peak_var    193002 non-null float64
3   price_peak_var        193002 non-null float64
4   price_mid_peak_var    193002 non-null float64
5   price_off_peak_fix    193002 non-null float64
6   price_peak_fix        193002 non-null float64
7   price_mid_peak_fix    193002 non-null float64
dtypes: float64(6), object(2)
memory usage: 11.8+ MB
```

In [10]:

```
client_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14606 entries, 0 to 14605
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    14606 non-null object
1   channel_sales         14606 non-null object
2   cons_12m              14606 non-null int64
3   cons_gas_12m          14606 non-null int64
4   cons_last_month       14606 non-null int64
5   date_activ            14606 non-null object
6   date_end              14606 non-null object
7   date_modif_prod       14606 non-null object
8   date_renewal          14606 non-null object
9   forecast_cons_12m     14606 non-null float64
10  forecast_cons_year     14606 non-null int64
11  forecast_discount_energy 14606 non-null float64
12  forecast_meter_rent_12m 14606 non-null float64
13  forecast_price_energy_off_peak 14606 non-null float64
14  forecast_price_energy_peak 14606 non-null float64
15  forecast_price_pow_off_peak 14606 non-null float64
16  has_gas               14606 non-null object
17  imp_cons              14606 non-null float64
18  margin_gross_pow_ele  14606 non-null float64
19  margin_net_pow_ele    14606 non-null float64
20  nb_prod_act           14606 non-null int64
21  net_margin            14606 non-null float64
22  num_years_antig       14606 non-null int64
23  origin_up             14606 non-null object
24  pow_max               14606 non-null float64
25  churn                 14606 non-null int64
dtypes: float64(11), int64(7), object(8)
memory usage: 2.9+ MB
```

Statistics Now let's look at some statistics about the datasets. We can do this by using the `describe()` method.

In [11]: `client_df.describe()`

Out[11]:

	cons_12m	cons_gas_12m	cons_last_month	forecast_cons_12m	forecast_cons_year	forecast
count	1.460600e+04	1.460600e+04	14606.000000	14606.000000	14606.000000	
mean	1.592203e+05	2.809238e+04	16090.269752	1868.614880	1399.762906	
std	5.734653e+05	1.629731e+05	64364.196422	2387.571531	3247.786255	
min	0.000000e+00	0.000000e+00	0.000000	0.000000	0.000000	
25%	5.674750e+03	0.000000e+00	0.000000	494.995000	0.000000	
50%	1.411550e+04	0.000000e+00	792.500000	1112.875000	314.000000	
75%	4.076375e+04	0.000000e+00	3383.000000	2401.790000	1745.750000	
max	6.207104e+06	4.154590e+06	771203.000000	82902.830000	175375.000000	

In [12]: `price_df.describe()`

Out[12]:

	price_off_peak_var	price_peak_var	price_mid_peak_var	price_off_peak_fix	price_peak_fix	price
count	193002.000000	193002.000000	193002.000000	193002.000000	193002.000000	
mean	0.141027	0.054630	0.030496	43.334477	10.622875	
std	0.025032	0.049924	0.036298	5.410297	12.841895	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.125976	0.000000	0.000000	40.728885	0.000000	
50%	0.146033	0.085483	0.000000	44.266930	0.000000	
75%	0.151635	0.101673	0.072558	44.444710	24.339581	
max	0.280700	0.229788	0.114102	59.444710	36.490692	

Data visualization

If you're working in Python, two of the most popular packages for visualization are matplotlib and seaborn. We highly recommend you use these, or at least be familiar with them because they are ubiquitous!

Below are some functions that you can use to get started with visualizations.

In [14]:

```
def plot_stacked_bars(dataframe, title_, size_=(18, 10), rot_=0, legend_="upper right")
    """
    Plot stacked bars with annotations
    """
    ax = dataframe.plot(
        kind="bar",
        stacked=True,
        figsize=size_,
        rot=rot_,
```

```

        title=title_)

    # Annotate bars
    annotate_stackedBars(ax, fontsize=14)
    # Rename Legend
    plt.legend(["Retention", "Churn"], loc=legend_)
    # Labels
    plt.ylabel("Company base (%)")
    plt.show()

def annotate_stackedBars(ax, pad=0.99, colour="white", fontsize=13):
    """
    Add value annotations to the bars
    """

    # Iterate over the plotted rectangles/bars
    for p in ax.patches:

        # Calculate annotation
        value = str(round(p.get_height(),1))
        # If value is 0 do not annotate
        if value == '0.0':
            continue
        ax.annotate(
            value,
            ((p.get_x()+ p.get_width()/2)*pad-0.05, (p.get_y()+p.get_height()/2)*pad),
            color=colour,
            size=fontsize)

def plot_distribution(dataframe, column, ax, bins=50):
    """
    Plot variable distribution in a stacked histogram of churned or retained company
    """

    # Create a temporal dataframe with the data to be plot
    temp = pd.DataFrame({"Retention": dataframe[dataframe["churn"]==0][column],
                        "Churn":dataframe[dataframe["churn"]==1][column]})
    # Plot the histogram
    temp[["Retention","Churn"]].plot(kind='hist', bins=bins, ax=ax, stacked=True)
    # X-axis Label
    ax.set_xlabel(column)
    # Change the x-axis to plain style
    ax.ticklabel_format(style='plain', axis='x')

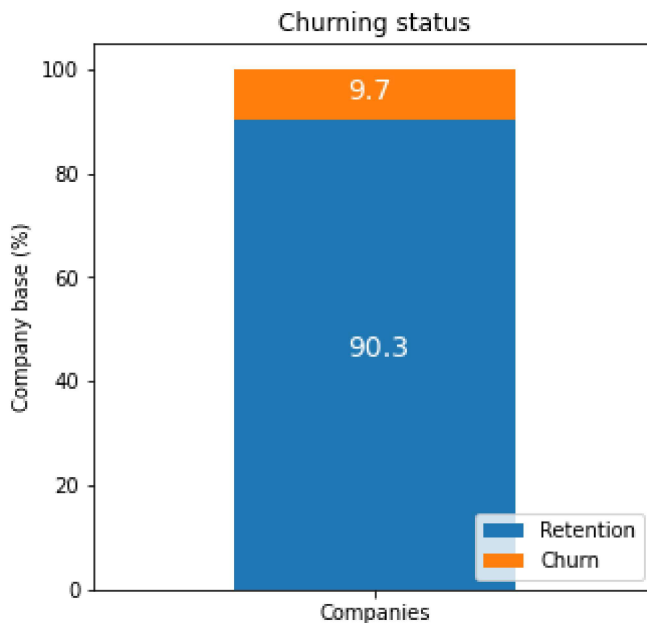
```

The first function plot_stacked_bars is used to plot a stacked bar chart. An example of how you could use this is shown below:

```


In [15]: churn = client_df[['id', 'churn']]
churn.columns = ['Companies', 'churn']
churn_total = churn.groupby(churn['churn']).count()
churn_percentage = churn_total / churn_total.sum() * 100
plot_stacked_bars(churn_percentage.transpose(), "Churning status", (5, 5), legend_=

```



The second function `annotate_bars` is used by the first function, but the third function `plot_distribution` helps you to plot the distribution of a numeric column. An example of how it can be used is given below:

```
In [16]: consumption = client_df[['id', 'cons_12m', 'cons_gas_12m', 'cons_last_month', 'imp_c
fig, axs = plt.subplots(nrows=1, figsize=(18, 5))
plot_distribution(consumption, 'cons_12m', axs)
```



The process of building a predictive model for customer churn based on client data and pricing data. Below are the steps you can follow, along with Python code snippets for each step. We'll be using Python and popular libraries like Pandas, Scikit-Learn, and Matplotlib for data processing, modeling, and visualization.

Step 1: Data Preprocessing Load and preprocess your data. This may involve handling missing values, encoding categorical variables, and splitting the data into training and testing sets.

```
In [ ]: # Import necessary libraries
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
```

```
# Load your client data and pricing data into a DataFrame
client_data = pd.read_csv('client_data.csv')
price_data = pd.read_csv('price_data.csv')

# Merge the datasets on a common key (e.g., customer ID)
merged_data = pd.merge(client_data, price_data, on='customer_id', how='inner')

# Encode categorical variables if needed (e.g., industry)
encoder = LabelEncoder()
merged_data['industry_encoded'] = encoder.fit_transform(merged_data['industry'])

# Split the data into training and testing sets
X = merged_data[['price', 'industry_encoded', 'contract_duration']]
y = merged_data['churn']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize numerical features (price and contract_duration)
scaler = StandardScaler()
X_train[['price', 'contract_duration']] = scaler.fit_transform(X_train[['price', 'contract_duration']])
X_test[['price', 'contract_duration']] = scaler.transform(X_test[['price', 'contract_duration']])

# Create and train a Logistic Regression model
model = LogisticRegression(random_state=42)
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')

# Generate a classification report
report = classification_report(y_test, y_pred)
print(report)

# Generate a confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Visualize the confusion matrix
plt.figure(figsize=(8, 6))
plt.imshow(conf_matrix, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xticks([0, 1], ['Not Churn', 'Churn'])
plt.yticks([0, 1], ['Not Churn', 'Churn'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

Name-Prashant Sharma email-iamprashant8008@gmail.com DATA Associate @NATIONAL AUSTRALIA BANK